



Experiment - 5

Student Name: Aryan Dhiman

UID: 23BCS12638

Branch: BE-CSE

Section/Group: KRG_3B

Semester: 5th

Date of Performance: 22/9/25

Subject Name: Advanced Database and Management System

Subject Code: 23CSP-333

1. Problem Description/Aim:

Medium-Problem Title: Generate 1 million records per ID in 'transaction_data' using generate_series() and random() ,create a normal view and a materialized view 'sales_summary' with aggregated metrics (total_quantity_sold , total_sales, total_orders) , and compare their performance and execution time.

Procedure (Step-by-Step):

1. Create a large dataset:
 - Create a table names transaction_data (id , value) with 1 million records.
 - take id 1 and 2, and for each id, generate 1 million records in value column
 - Use Generate_series () and random() to populate the data.
2. Create a normal view and materialized view to for sales_summary, which includes total_quantity_sold, total_sales, and total_orders with aggregation.
3. Compare the performance and execution time of both.

Sample Output Description:

The transaction_data table has 2 million rows (1 million per ID) with random values. The normal view sales_summary computes aggregates on the fly, while the materialized view sales_summary_mv stores precomputed results. Queries on the materialized view are much faster, but it needs refreshing when data changes, whereas the normal view always shows up-to-date results.

Hard-Problem Title: Create restricted views in the sales database to provide summarized, non-sensitive data to the reporting team, and control access using DCL commands(GRANT and REVOKE).

Procedure (Step-by-Step):

1. Create restricted views-
 - Define views that show only **aggregated sales data** (e.g., total_sales, total_orders) without exposing sensitive columns like customer details or payment info.

2. Assign access to reporting team(or client)-
 - Use “GRANT SELECT ON view_name TO reporting_user;” to give access.
3. Revoke access if needed.
 - Use “REVOKE SELECT ON view_name FROM reporting_user;” to remove access.
4. Verify access
 - Reporting users can query the view but cannot access base tables directly, ensuring security.

Sample Output Description:

The result shows the restricted view providing summarized sales data only like

- Columns shown are - product_id,total_quantity_sold, total_sales, total_orders
- Columns hidden are - Customer names, addresses, payment details

A reporting user querying the view sees something like :

- Product 101 - 5000 units sold, total sales Rs. 12,50,000,500 orders.
- Product 102 - 3200 units sold, total sales Rs. 8,60,000,320 orders.

When the user tries to query the base “sales_transactions” table directly, access is denied, enforcing security.

- 2. Objective:** To design and implement secure, efficient data access mechanisms by creating large-scale transaction datasets, summarizing them through normal and materialized views for performance comparison, and enforcing restricted access to sensitive data using views and DCL commands.

3. SQL QUERY AND OUTPUTS -

-----MEDIUM LEVEL PROBLEM-----

Create table TRANSACTION_DATA(id int,val decimal);

INSERT INTO TRANSACTION_DATA(ID,VAL)

SELECT 1,RANDOM()

FROM GENERATE_SERIES(1,1000000);

INSERT INTO TRANSACTION_DATA(ID,VAL)

SELECT 2,RANDOM()

FROM GENERATE_SERIES(1,1000000);

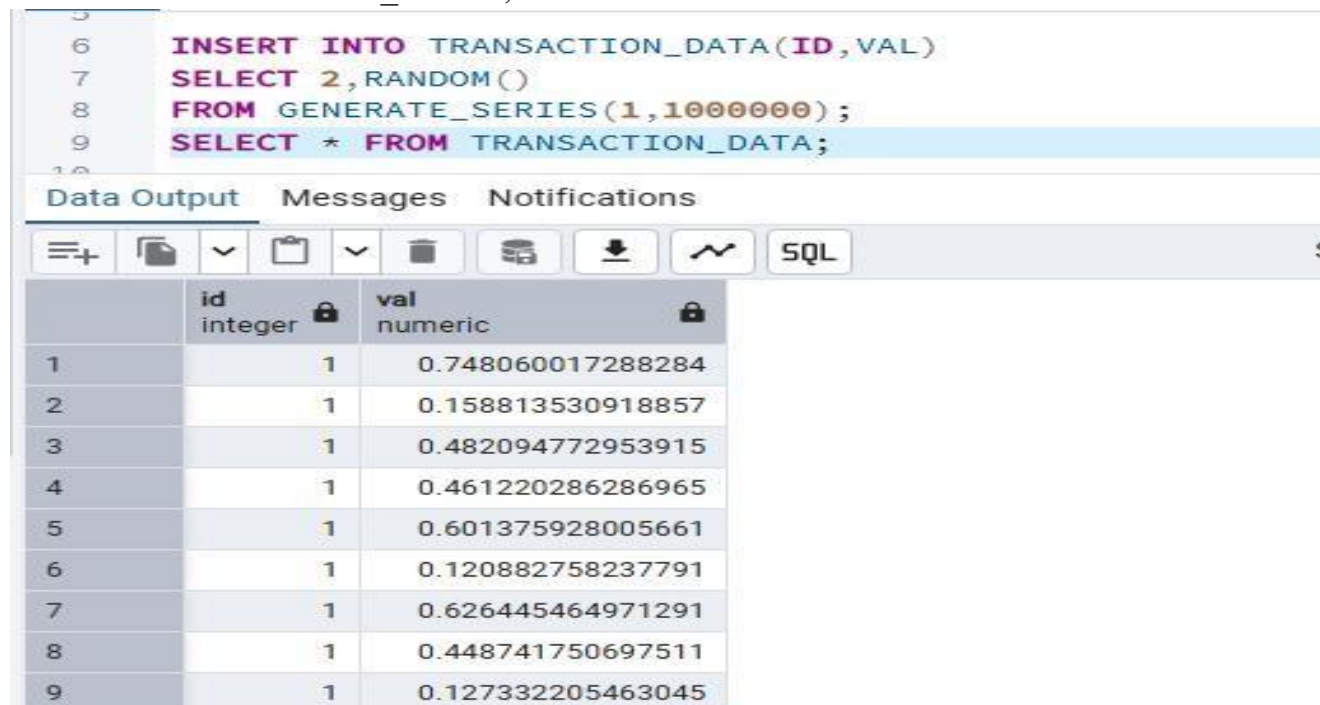
SELECT * FROM TRANSACTION_DATA;

```
CREATE or REPLACE VIEW SALES_SUMMARY AS  
SELECT  
ID,  
COUNT(*) AS total_quantity_sold,  
sum(val) AS total_sales,  
count(distinct id) AS total_orders  
FROM TRANSACTION_DATA  
GROUP BY ID;
```

```
EXPLAIN ANALYZE  
SELECT * FROM SALES_SUMMARY;
```

```
CREATE MATERIALIZED VIEW SALES_SUMM AS  
SELECT  
ID,  
COUNT(*) AS total_quantity_sold,  
sum(val) AS total_sales,  
count(distinct id) AS total_orders  
FROM TRANSACTION_DATA  
GROUP BY ID;
```

```
EXPLAIN ANALYZE  
SELECT * FROM SALES_SUMM;
```



The screenshot shows a SQL IDE interface. The top pane displays a query that inserts random data into the TRANSACTION_DATA table and then selects all records. The bottom pane shows the 'Data Output' tab with a table of 9 rows. The table has two columns: 'id' (integer) and 'val' (numeric). The 'id' column contains values from 1 to 9, and the 'val' column contains random decimal values.

| | id integer | val numeric |
|---|---------------|-------------------|
| 1 | 1 | 0.748060017288284 |
| 2 | 1 | 0.158813530918857 |
| 3 | 1 | 0.482094772953915 |
| 4 | 1 | 0.461220286286965 |
| 5 | 1 | 0.601375928005661 |
| 6 | 1 | 0.120882758237791 |
| 7 | 1 | 0.626445464971291 |
| 8 | 1 | 0.448741750697511 |
| 9 | 1 | 0.127332205463045 |

```
21 SELECT * FROM SALES_SUMMARY; /*Simple view */
```

Data Output Messages Notifications

| | id integer | total_quantity_sold bigint | total_sales numeric | total_orders bigint |
|---|---------------|-------------------------------|-------------------------------|------------------------|
| 1 | 1 | 2000000 | 1000226.201610874170319933640 | 1 |
| 2 | 2 | 1000000 | 499473.47586932728250459408 | 1 |

```
20 EXPLAIN ANALYZE
21 SELECT * FROM SALES_SUMMARY; /*Simple view */
```

Data Output Messages Notifications

| | QUERY PLAN text |
|---|---|
| 1 | GroupAggregate (cost=471514.97..509014.99 rows=2 width=52) (a |
| 2 | Group Key: transaction_data.id |
| 3 | -> Sort (cost=471514.97..479014.97 rows=3000000 width=15) (ac |
| 4 | Sort Key: transaction_data.id |
| 5 | Sort Method: external merge Disk: 73504kB |
| 6 | -> Seq Scan on transaction_data (cost=0.00..46224.00 rows=3 |
| 7 | Planning Time: 0.135 ms |
| 8 | Execution Time: 4396.880 ms |

```
33 SELECT * FROM SALES_SUMM; /*Materialized view*/
```

Data Output Messages Notifications

| | id integer | total_quantity_sold bigint | total_sales numeric | total_orders bigint |
|---|---------------|-------------------------------|------------------------------|------------------------|
| 1 | 1 | 1000000 | 500106.667545326356598143529 | 1 |
| 2 | 2 | 1000000 | 499473.47586932728250459408 | 1 |

| | |
|--|---|
| 32 | EXPLAIN ANALYZE |
| 33 | SELECT * FROM SALES_SUMM; /*Materialized view*/ |
| Data Output Messages Notifications | |
| <div> <div> <div>+</div> <div>SQL</div> </div> <div>Showing rows: 1</div> </div> | |
| | <div> <div>QUERY PLAN</div> <div>text</div> <div>lock</div> </div> |
| 1 | Seq Scan on sales_summ (cost=0.00..20.20 rows=1020 width=52) (actual time=0.017..0.018 rows=2 loops=... |
| 2 | Planning Time: 0.063 ms |
| 3 | Execution Time: 0.032 ms |

OUTPUT -

As we can see that the execution time using the materialized view is very less as compared to the simple view's execution time.

-----HARD PROBLEM -----

```
CREATE TABLE customer_data (
  transaction_id SERIAL PRIMARY KEY,
  customer_name VARCHAR(100),
  email VARCHAR(100),
  phone VARCHAR(15),
  payment_info VARCHAR(50), -- sensitive
  order_value DECIMAL,
  order_date DATE DEFAULT CURRENT_DATE
);
```

-- Insert sample data

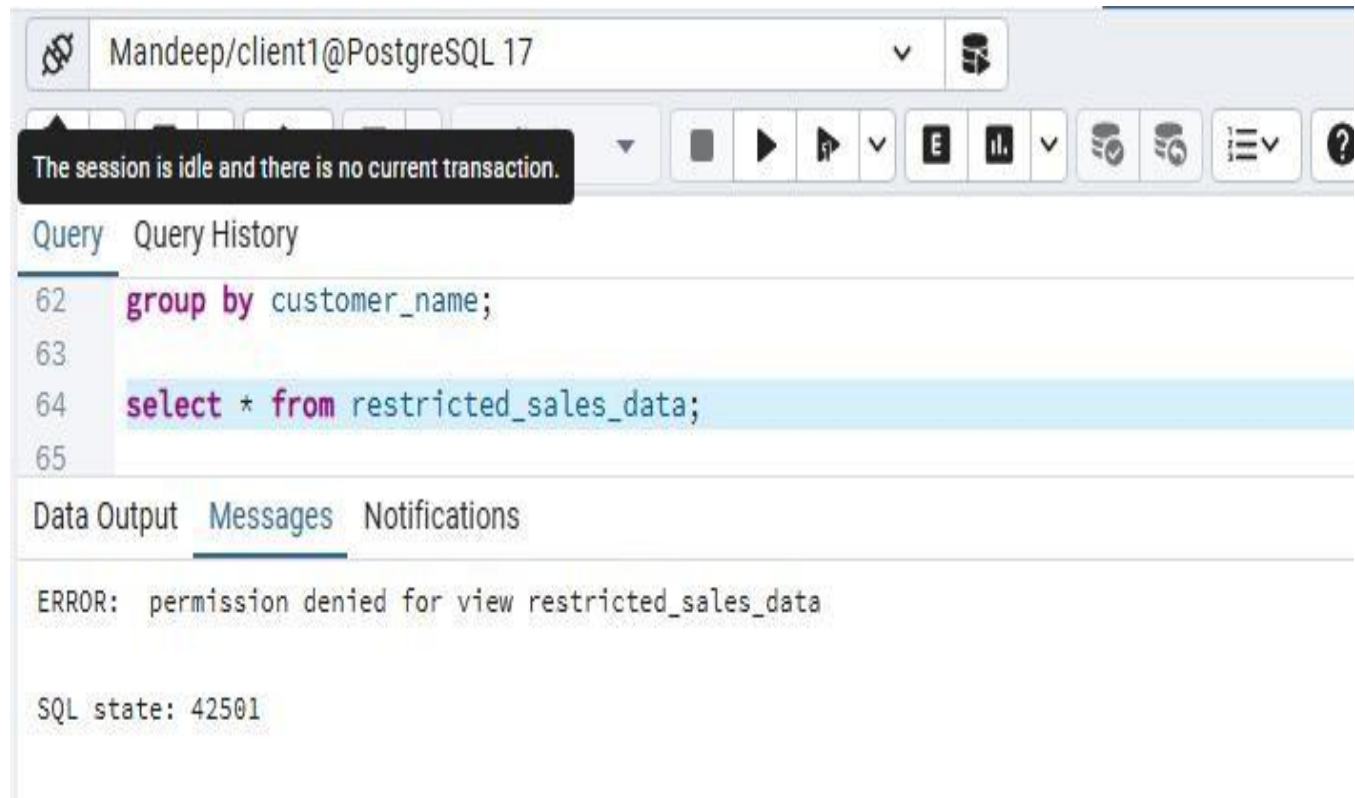
```
INSERT INTO customer_data (customer_name, email, phone, payment_info, order_value)
VALUES
('Akshara Chauhan', 'mandeep@example.com', '9040122324', '1234-5678-9012-3456', 500),
('Akshara Chauhan', 'mandeep@example.com', '9040122324', '1234-5678-9012-3456', 1000),
('Ishika Chauhan', 'jaskaran@example.com', '9876543210', '9876-5432-1098-7654', 700),
('Akshara Chauhan', 'jaskaran@example.com', '9876543210', '9876-5432-1098-7654', 300);
```

```
CREATE OR REPLACE VIEW RESTRICTED_SALES_DATA AS  
SELECT
```

```
CUSTOMER_NAME,  
COUNT(*) AS total_orders,  
SUM(order_value) as total_sales  
from customer_data  
group by customer_name;
```

```
select * from restricted_sales_data;
```

```
CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';  
GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;  
REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```



The screenshot shows a PostgreSQL client interface with the connection string 'Mandeep/client1@PostgreSQL 17'. A message box states 'The session is idle and there is no current transaction.' The query editor shows two SQL statements: 'group by customer_name;' and 'select * from restricted_sales_data;'. The second statement is highlighted in blue. Below the query editor, the 'Messages' tab is active, displaying an error: 'ERROR: permission denied for view restricted_sales_data' with the SQL state '42501'.

Mandeep/client1@PostgreSQL 17

The session is idle and there is no current transaction.

Query Query History

```
62 group by customer_name;  
63  
64 select * from restricted_sales_data;  
65
```

Data Output Messages Notifications

ERROR: permission denied for view restricted_sales_data

SQL state: 42501

Mandeep/postgres@PostgreSQL 17

Query Query History

```
65  
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';  
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;  
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

GRANT

Query returned successfully in 154 msec.

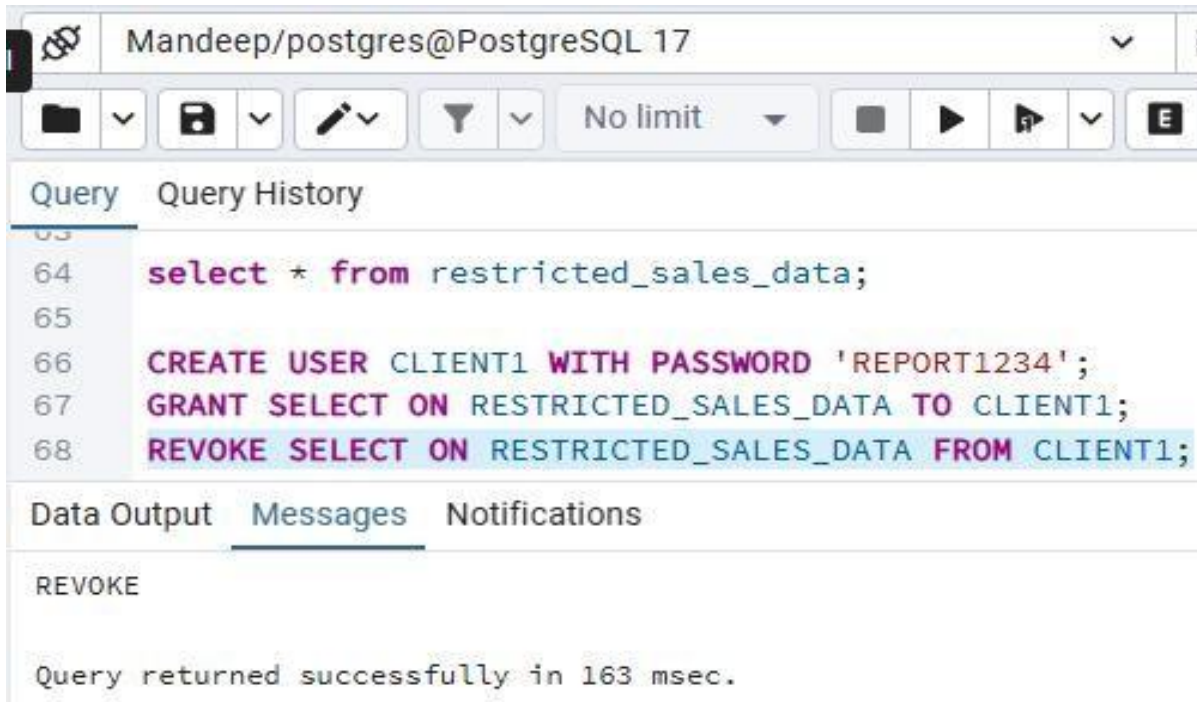
Mandeep/client1@PostgreSQL 17

Query Query History

```
62 group by customer_name;  
63  
64 select * from restricted_sales_data;  
65
```

Data Output Messages Notifications

| | customer_name character varying (100) | total_orders bigint | total_sales numeric |
|---|--|------------------------|------------------------|
| 1 | Jaskaran Singh | 2 | 1000 |
| 2 | Mandeep Kaur | 2 | 1500 |



Mandeep/postgres@PostgreSQL 17

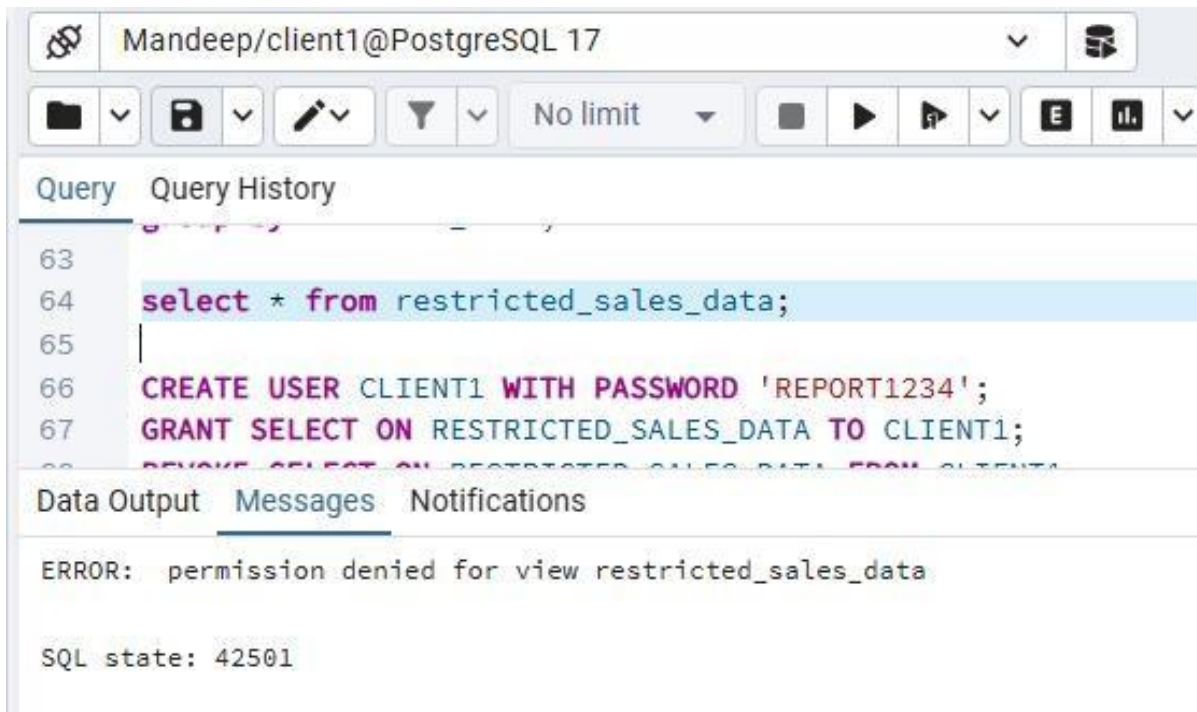
Query Query History

```
63  
64 select * from restricted_sales_data;  
65  
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';  
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;  
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

REVOKE

Query returned successfully in 163 msec.



Mandeep/client1@PostgreSQL 17

Query Query History

```
63  
64 select * from restricted_sales_data;  
65  
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';  
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;  
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

ERROR: permission denied for view restricted_sales_data

SQL state: 42501